

Parallel Computing

Academic year – 2020/21, spring semester
Computer science

Lecture 7

Lecturer, instructor:

Balakshin Pavel Valerievich

(pvbalakshin@itmo.ru; pvbalakshin@hdu.edu.cn)

Assistant:

Liang Tingting

(liangtt@hdu.edu.cn)

Parallel computing paradigm

- **Explicit locking** for synchronization when addressing to the shared data.
- **Non-blocking algorithms** (lockless, lockfree algorithms) – use elementary atomic operations only.
- **Software Transactional Memory** – non-synchronous use of divisible data, followed by eliminating conflict.

Explicit locking

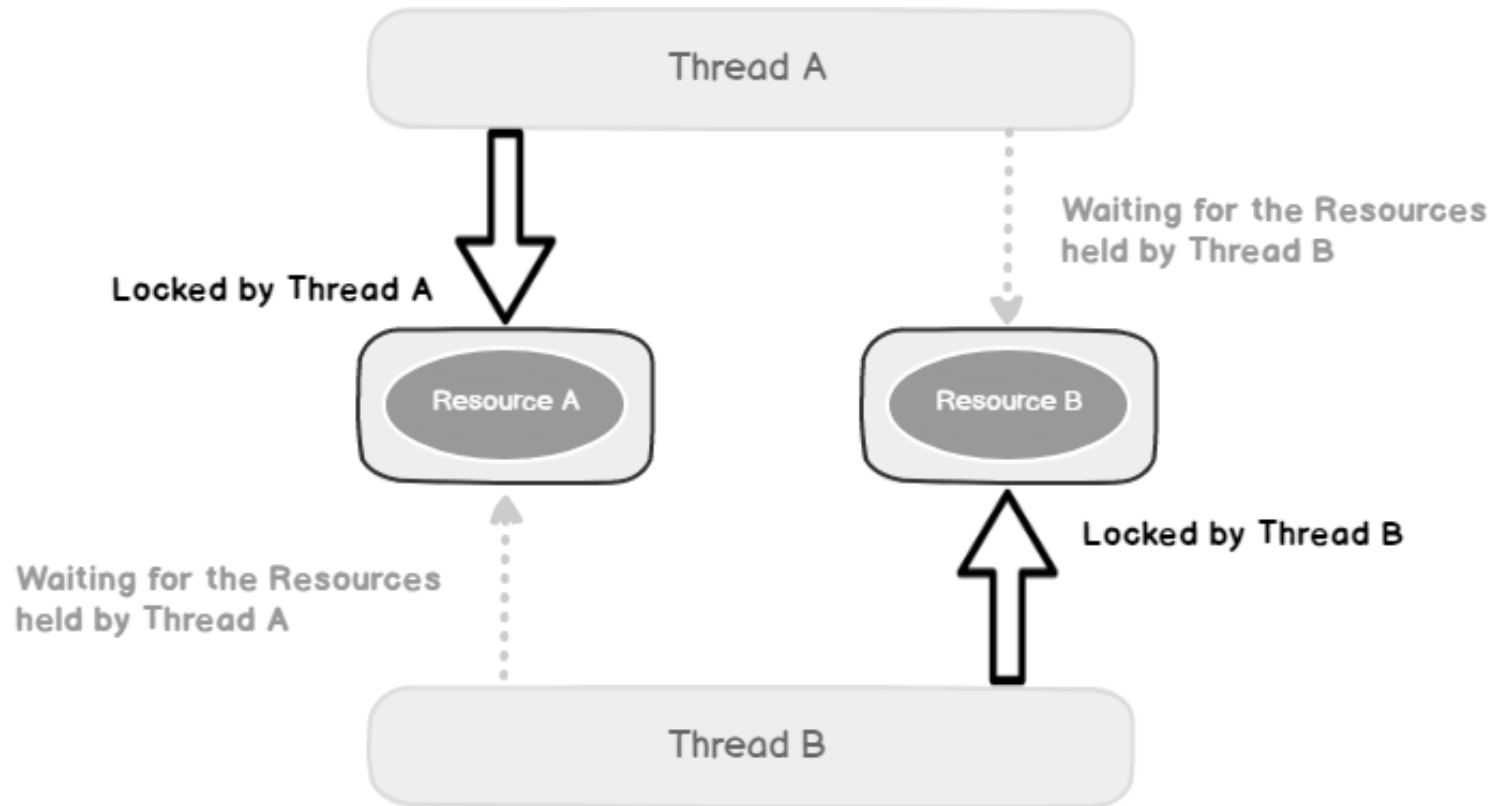
1. The most popular one:

- *Semaphore* – an object that limits the number of threads that can enter this area of code.
- *Reader/writer semaphore* gives the threads only read or write permissions, and while writing data to one thread, the rest of the threads do not have access to the resource.
- *Mutex* is a special case of a semaphore, in which only one thread can capture a given area of code.
- *Spinlock* is a lock in which a thread in a loop waits for a resource to free.
- *Seqlock* – synchronization mechanism designed to quickly record a variable in multiple threads.
- *Barrier* – a part of code in which the state of threads is synchronized.

Explicit locking

2. Unstable waiting time when locked without upper limit.
3. Deadlocks are possible.
4. Construction/removal overhead and waiting costs.
5. Difficulties in debugging.
6. Poor scalability.
7. Priority inversion is possible.

Deadlock



According to G. Gorelkin "Deadlocks, Livelocks и Starvation"


Deadlock (2)

```
omp_lock_t lock1, lock2;
omp_init_lock(&lock1); /* initialization of lock 1 */
omp_init_lock(&lock2); /* initialization of lock 2 */
#pragma omp parallel sections
{
    #pragma omp section
    {
        omp_set_lock(&lock1);    omp_set_lock(&lock2);
        omp_unset_lock(&lock2);  omp_unset_lock(&lock1);
    }
    #pragma omp section
    {
        omp_set_lock(&lock2);    omp_set_lock(&lock1);
        omp_unset_lock(&lock1);  omp_unset_lock(&lock2);
    }
}
```

Non-blocking algorithms

```
1  /* Pseudocode of the instruction returning boolean value in C syntax */
2  int cas( int* addr, int old, int new )
3  {
4      if ( *addr != old )
5          return 0;
6
7      *addr = new;
8      return 1;
9  }
```

1. Block resource.
2. Read value of the variable.
3. Process some computations.
4. Write new value of the variable.
5. Release block of resource.

- 
1. Read value of the variable.
 2. Process some computations.
 3. Process cmpxchg operation with new value of variable assuming that old value is still the same.
 4. Repeat steps 2-3 if value of the variable was changed by another thread.

Non-blocking algorithms

- No standard limitations of blocking algorithms.
- Quite high input level.
- Program became bigger and more difficult to maintain and support.
- The limited number of atomic operations leads to a significant complication of the algorithm.

Transactional memory

- Easy of use – enclosing code sections in a transaction block.
- Absence of blocking.
- Working good and fast in case of small amount of access conflicts (can be roll back in case of failure).
- Ideally it requires hardware support (execution of transactions not by compiler, but processor with special internal registers to store old values), which is not yet implemented in all processors.
- Supported in GCC since version 4.7:
 - `_transaction atomic {...}` – an indication that the code block is a transaction;
 - `_transaction_relaxed {...}` – an indication that the unsafe code inside the block does not lead to side effects;
 - `_transaction cancel` – explicit transaction cancellation.

Transactional memory (2)

```
int contains(int value)
{
    int result;
    node_t *prev, *next;

    __transaction_atomic {
        prev = set->head;
        next = prev->next;
        while (next->val < val) {
            prev = next;
            next = prev->next;
        }
        result = (next->val == val);
    }
    return result;
}
```

```
int add(int value)
{
    int result;
    node_t *prev, *next;

    __transaction_atomic {
        prev = set->head;
        next = prev->next;
        while (next->val < val) {
            prev = next;
            next = prev->next;
        }
        result = (next->val != val);
        if (result) {
            prev->next = new_node(val, next);
        }
    }
    return result;
}
```

Pthreads

Pthreads - the POSIX (IEEE) standard, which defines the API for thread management.

Key Pthreads tools provide:

- Create, delete, merge threads.
- Shared access to shared variables.
- Synchronization of threads.


Compiling Pthreads-program

Intel GNU/Linux	<code>icc -pthread</code>
	<code>icpc -pthread</code>
PGI GNU/Linux	<code>pgcc -lpthread</code>
	<code>pgCC -lpthread</code>
GNU GCC GNU/Linux, Blue Gene	<code>gcc -pthread</code>
	<code>g++ -pthread</code>
IBM Blue Gene	<code>bgxlc_r / bgcc_r</code>
	<code>bgxlc_r, bgxlc++_r</code>

There are no "native" Microsoft tools for compiling for Windows, but you can use third-party libraries.

Creation and deletion of threads

```
int main(int argc, char *argv[])
{
    int a = 1, b = 2;
    #pragma omp parallel sections
    #pragma omp parallel section
    function_1(a);
    #pragma omp parallel section
    function_1(b);
}
```



```
int main(int argc, char *argv[])
{
    pthread_t t1, t2;
    int a = 1, b = 2;
    pthread_create(&t1, NULL, function_1, (void*)&a);
    pthread_create(&t2, NULL, function_1, (void*)&b);
    pthread_exit(NULL);
}
```

Creation and deletion of threads (2)

```
#include <pthread.h>
#define NTHREADS 5

void *thread_fun(void *threadid) {
    long tid = (long)threadid;
    printf("Hello from %ld!\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[]) {
    pthread_t threads[NTHREADS];
    int rc; long t;

    for (t = 0; t < NTHREADS; t++) {
        rc = pthread_create(&threads[t], NULL, thread_fun, (void *)t);
        if (rc) {
            printf("ERROR %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

Creation and deletion of threads (3)

```
#include <pthread.h>
#define NTHREADS 5

void *thread_fun(void *threadid) {
    long tid = (long)threadid;
    printf("Hello from %ld!\n", tid);
    pthread_exit(NULL);
}

int main(int argc, char *argv[]) {
    pthread_t threads[NTHREADS];
    int rc; long t;

    for (t = 0; t < NTHREADS; t++) {
        rc = pthread_create(&threads[t], NULL, thread_fun, (void *)t);
        if (rc) {
            printf("ERROR %d\n", rc);
            exit(-1);
        }
    }
    pthread_exit(NULL);
}
```

```
int main(int argc, char *argv[])
{
    #pragma omp parallel num_threads(5)
    printf("Hello from %u\n!",
        omp_get_thread_num());
}
```

Creation and deletion of threads (4)

```
$ gcc -pthread -o prog ./prog.c  
$ ./prog  
Hello from 1!  
Hello from 4!  
Hello from 0!  
Hello from 2!  
Hello from 3!
```


Threadpool consolidation

```
#include <pthread.h>
#define NTHREADS 5

// ...

int main(int argc, char *argv[]) {
    pthread_t threads[NTHREADS];
    int rc; long t;
    void *status;

    for (t = 0; t < NTHREADS; t++) {
        rc = pthread_create(&threads[t], NULL, thread_fun,
                           (void *)t);
    }

    for (t = 0; t < NTHREADS; t++) {
        rc = pthread_join(threads[t], &status);
    }
    pthread_exit(NULL);
}
```

Create threads and wait for all threads in the main function

```
intervals = atoi(argv[1]);
numThreads = atoi(argv[2]);
threads = malloc(numThreads * sizeof(pthread_t));
threadID = malloc(numThreads * sizeof(int));
pthread_mutex_init(&piLock, NULL);
for (i = 0; i < numThreads; i++) {
    threadID[i] = i;
    pthread_create(&threads[i], NULL, computePI, threadID + i);
}
for (i = 0; i < numThreads; i++)
    pthread_join(threads[i], &retval);

printf("Estimation of pi is %32.30Lf \n", pi);
```

Shared variable protection

```
pthread_mutex_t piLock;
long double intervals;
int numThreads;

void *computePI(void *id)
{
    long double x, width, localSum = 0;
    int i, threadID = *((int*)id);    width = 1.0 / intervals;
    for (i = threadID ; i < intervals; i += numThreads) {
        x = (i + 0.5) * width;
        localSum += 4.0 / (1.0 + x * x);
    }
    localSum *= width;
    pthread_mutex_lock(&piLock);
    pi += localSum;
    pthread_mutex_unlock(&piLock);
    return NULL;
}
```

Threads synchronization without their removal (stop)

```
void * entry_point(void *arg)
{
    int rank = (int)arg;
    for(int row = rank * ROWS / THREADS; row < (rank + 1) * THREADS; ++row)
        for(int col = 0; col < COLS; ++col)
            DotProduct(row, col, initial_matrix, final_matrix);

    // Synchronization point
    int rc = pthread_barrier_wait(&barr);
    if(rc != 0 && rc != PTHREAD_BARRIER_SERIAL_THREAD)
    {
        printf("Could not wait on barrier\n");
        exit(-1);
    }

    for(int row = rank * ROWS / THREADS; row < (rank + 1) * THREADS; ++row)
        for(int col = 0; col < COLS; ++col)
            DotProduct(row, col, final_matrix, initial_matrix);
}

int main(int argc, char **argv)
{
    pthread_t thr[THREADS];

    // Barrier initialization
    if(pthread_barrier_init(&barr, NULL, THREADS))
```